

Supplementary Material for PrepBench: A Factorized Benchmark for Cost-Aware Data-Preparation Pipeline Search [Experiment, Analysis & Benchmark]

Jing Chang, Chang Liu, Jinbin Huang, Yu-Xuan Qiu, Shuyuan Zheng, Rui Mao, Jianbin Qin

This document accompanies the main paper and provides the artifact description, the released file schemas, the exact reproduction procedure, and the hardware/software environment. Every claim in the paper is reproducible from the released files described here. The artifact is available at:

<https://github.com/prepbench2027/prepbench>

1 ARTIFACT OVERVIEW AND CLAIMS SUPPORTED

The artifact is organized around complete run summaries and candidate-level histories, so that every number, table, and figure in the paper can be regenerated offline from released CSVs without re-running the expensive benchmark. The benchmark evaluates nine configurations (seven factorial combinations of the structure/-cost/evaluation/budget axes plus two random controls) across 56 OpenML datasets with five seeds each, at wall-clock budgets of 60 s, 300 s, and 600 s, totalling 2,520 runs per budget. Concretely, the artifact supports the following claims:

- **Mean scores, ranks, and omnibus tests** for the controlled 56-dataset matrix, including the critical-difference diagram, are regenerated from `experiment_records.csv`.
- **Pairwise Wilcoxon and TOST equivalence tests** are regenerated from `experiment_records.csv` and the derived paired per-dataset score tables.
- **Zero-penalty sensitivity** is regenerated from `results/sensitivity/` $\hookrightarrow$ and supports the claim that the random-control ranking is not explained by the default estimated-cost or graph-complexity penalties.
- **Stronger guided-EA sensitivity** is regenerated from `results/` $\hookrightarrow$ `sensitivity/strong_ea300/` and supports the scoped claim that optimizer strength matters while the combined cost-aware variant remains below the plain DAG in final accuracy.
- **Random-space fairness audit** is regenerated from `random_fairness_audit` $\hookrightarrow$ `.csv` and records candidate-level structure, workflow uniqueness, and runtime diagnostics for the zero-penalty runs.
- **Paired budget deltas** and the budget-scaling figure are regenerated from `experiment_records.csv`.
- **Evaluation-allocation counters** are regenerated from `experiment_records` $\hookrightarrow$ `.csv`.
- **Staged-screening diagnostics** and the calibration scatter plot are regenerated from `search_history.csv`.
- **Cost-model calibration**, including the static-cost vs. runtime scatter plot moved out of the main paper for space, is regenerated from `search_history.csv`.
- **Incumbent convergence** is regenerated from `search_history.csv`.
- **Efficiency frontier** is regenerated from `experiment_records.csv`.

- **Provenance and robustness**, including the completed-evaluation-only block summarized in the main paper, come from coverage- and trace-audit analyses over `experiment_records.csv` and `search_history` $\hookrightarrow$ `.csv`.
- **External CtxPipe reference** is regenerated from `ctxpipe_scores` $\hookrightarrow$ `csv` and `ctxpipe_paired_comparison.csv`; it is contextual evidence, not a controlled ablation.

2 REPOSITORY LAYOUT

```
prepbench-vldb-eab-artifact/  
  README.md  
  requirements.txt  
  paper/  
    paper-vldb.pdf  
    supplementary_prepbench.pdf  
  code/  
    baselines/  
    parallel_evolutionary_pipeline/  
    data_loader/  
    experiments/  
      run_matrix.py  
      analyze_icde_eab_extras.py  
      summarize_sensitivity_runs.py  
      make_paper_result_bundle.py  
    figure_scripts/  
    scripts/  
  results/  
    main_matrix/  
      eab60/ eab300/ eab600/  
        experiment_records.csv  
        run_metrics.csv  
        search_history.csv  
        run_manifest.json  
        paper_result_bundle/  
    sensitivity/  
      zpen300/ zpen600/  
      strong_ea300/  
      zero_penalty_audit/  
        random_fairness_audit.csv  
      strongEA_compare/  
      sensitivity_summary_by_config.csv  
      sensitivity_paired_tests.csv  
    external_baselines/  
      ctxpipe_scores.csv  
      ctxpipe_paired_comparison.csv  
  figures/  
    data/  
      manifest.csv
```

3 RELEASED FILE SCHEMAS

Every released CSV is documented below; column names match the analysis scripts exactly.

`experiment_records.csv` (*per budget*). One row per (dataset, seed, configuration, budget) run — the unit used for all score-based statistics:

dataset, config, seed, budget, score,
score_source in {full_candidate, stage1, fallback},

estimated_cost, serial_cost, critical_path_cost,
runtime_seconds, branches, nodes, depth,
n_full_evals, n_screened_out, n_pruned

run_metrics.csv (*per run directory*). Augmented per-run metrics used for anytime, efficiency, and allocation analyses:

dataset, config, seed, budget, score,
n_full_evals, n_ok_evals, n_screened_out, n_pruned,
branches, nodes, depth, time_to_target, anytime_auc

search_history.csv (*per run directory*). One row per candidate event (the trace used for allocation, screening, and cost-calibration diagnostics):

dataset, config, seed, budget, candidate_index,
status in {ok, screened_out, pruned, timeout, failed, fallback},
score_source, score, stage1_score, full_score,
best_so_far, improved_incumbent,
estimated_serial_cost, estimated_critical_path_cost,
elapsed_seconds, branches, nodes, operator, generation

summary_by_config.csv. Per-config aggregate statistics across datasets and seeds:

budget, config, mean_score, median_score, std_score,
mean_evals, mean_ok_evals, mean_screened_out, mean_pruned,
mean_branches, mean_nodes, mean_depth,
mean_time_to_target, mean_anytime_auc, n_runs

run_manifest.json. Experiment provenance: budget, list of seeds, list of configs, list of datasets, penalty settings, population/offspring settings when changed, and the run tag used to separate main-matrix and revision-sensitivity runs.

sensitivity/ files. Sensitivity outputs support the optimizer/objective checks added to the main parser:

- sensitivity_summary_by_config.csv: aggregate means, medians, ranks, and run counts for zero-penalty and stronger-EA runs.
- sensitivity_paired_tests.csv: paired per-dataset Wilcoxon comparisons after averaging seeds.
- random_fairness_audit.csv: candidate-level audit of workflow uniqueness, mean nodes, mean branches, runtime, success rate, and entropy diagnostics for the random and guided DAG variants.
- Per-run experiment_records.csv, run_metrics.csv, and search_history. $\hookrightarrow$ csv under each sensitivity run directory.

trace_summary/ files.

- status_summary.csv: distinct status values $\times$ config $\times$ budget.
- per_run_incumbent.csv: one row per (dataset, config, seed, budget) recording incumbent score, score_source, status counts, and evaluation counts.
- staged_metrics.csv: per (config, budget) Pearson and Spearman correlation between stage-1 and full score, promotion rate, and incumbent-threat rate.
- stage_full_sample.csv: reservoir sample of (stage1_score, full_score) pairs for scatter plotting (Fig. 4).

paper_result_bundle/ files. Generated by make_paper_result_bundle.py:

- stat_tests.csv: Friedman omnibus test and Wilcoxon signed-rank pairwise tests (each configuration vs. the DAG baseline D-S-F-A, and the random controls vs. the structured configurations).

- tost_sensitivity.csv: mean/median differences and within-margin flags for $\delta \in \{0.005, 0.01, 0.02\}$ against D-S-F-A.
- ctxpipe_scores.csv and ctxpipe_paired_comparison.csv: CtxPipe per-dataset reference scores and paired comparison against the controlled references.
- coverage_audit.csv: per-budget expected vs. observed runs at the run level and the candidate-trace level.
- trace_audit.csv: cross-validation between run_metrics and search_history $\hookrightarrow$ key sets.
- dataset_config_scores.csv, dataset_config_seed_scores.csv: per-dataset aggregated scores.

manifest.csv. The manifest has one row per dataset and records the canonical dataset name, the artifact-relative CSV path, the target column, the expected info.json path, the original DeepLine source URL, and the dataset statistics used in Table 2. The OpenML-id fields are retained for compatibility with external-baseline scripts but are left blank when the artifact uses the DeepLine CSV source directly; in that case csv_path and target are the authoritative loading fields.

Provenance contract. For every run with at least one completed full evaluation, the reported score in experiment_records.csv equals the score of the best incumbent in search_history.csv. The run-level allocation counters (n_full_evals, n_screened_out, n_pruned) match the aggregated candidate-status counts. Runs whose reported score has score_source = fallback (i.e., no full candidate evaluation completed) are excluded from the completed-evaluation-only robustness block.

4 REPRODUCTION PROCEDURE

The released results reproduce all paper numbers in minutes on a laptop; the expensive end-to-end benchmark is optional and supports resumption.

```
# From the artifact root:
cd prepbench-vldb-eab-artifact
conda create -n prepbench python=3.12
conda activate prepbench
pip install -r requirements.txt
```

```
python code/experiments/make_paper_result_bundle.py \
--results results/main_matrix \
--out /tmp/prepbench_result_check
```

```
python code/figure_scripts/cd_diagram_600.py
python code/figure_scripts/incumbent_convergence_600.py
python code/figure_scripts/budget_delta_score_vs_evals.py
python code/figure_scripts/staged_calibration_dctb_600.py
python code/figure_scripts/efficiency_frontier_600.py
python code/figure_scripts/cost_runtime_calibration.py
```

```
# Regenerate revision-sensitivity summaries
python code/experiments/summarize_sensitivity_runs.py \
--root results/sensitivity \
--out /tmp/prepbench_sensitivity_check
```

These scripts read only the released CSVs under results/ and write PDFs to figures/. They are deterministic and produce the figure panels used in the paper and the supplementary diagnostic plots.

```
# Re-run the full benchmark (optional)
python code/experiments/run_matrix.py \
--budget 600 --seeds 42 114 256 512 768 \
--out results/main_matrix/eab600 \
--processes 32 --per-eval-timeout 10 \
--run-tag main_600
```

Table 1: Environment fingerprint.

Field	Value
CPU	AMD EPYC 7543 32-Core Processor (dual socket)
Logical threads	128 (SMT)
Memory	512 GB
OS	Ubuntu 20.04.6 LTS
Parallel backend	joblib (loky)
Python	3.12.9
scikit-learn	1.6.0
numpy	1.26.4
scipy	1.16.3
pandas	2.3.3
dataclasses	0.6
matplotlib	3.10.1
python-dotenv	1.2.1

```
# Re-run a zero-penalty sensitivity block
python code/experiments/run_matrix.py \
  --budget 600 --seeds 42 114 256 512 768 \
  --configs D-Rand D-S-F-A D-C-T-B \
  --out results/sensitivity/zpen600 \
  --processes 64 --per-eval-timeout 10 \
  --fitness-cost-penalty 0 \
  --fitness-complexity-penalty 0 \
  --run-tag zero_penalty_600
```

```
# Re-run the stronger guided-EA sensitivity block
python code/experiments/run_matrix.py \
  --budget 300 --seeds 42 114 256 512 768 \
  --configs D-S-F-A D-C-T-B \
  --out results/sensitivity/strong_ea300 \
  --processes 64 --per-eval-timeout 10 \
  --fitness-cost-penalty 0 \
  --fitness-complexity-penalty 0 \
  --population-size 40 \
  --offspring-size 40 \
  --run-tag strong_ea300
```

This reproduces the raw runs and writes fresh `experiment_records.csv`, `run_metrics.csv`, and `search_history.csv`.

5 STATISTICAL PROTOCOL DETAILS

For completeness we record the exact procedures used in the paper.

- **Aggregation.** Scores are averaged over the five seeds within each (dataset, configuration) cell, giving a complete 56×9 block per budget; the dataset is the unit of analysis.
- **Omnibus + post-hoc.** A Friedman test per budget; on rejection, the Nemenyi post-hoc test with $k = 9$, $N = 56$, $q_{0.05} = 3.102$, giving $CD = 1.605$.
- **Pairwise.** Wilcoxon signed-rank on paired per-dataset scores, with Holm–Bonferroni family-wise control within each budget.
- **Equivalence.** Two one-sided tests (TOST) on paired per-dataset differences. The main margin is $\delta = 0.01$; the sensitivity analysis sweeps $\delta \in \{0.005, 0.01, 0.02\}$ in the equivalence-margin sensitivity table. Two configurations are declared equivalent only if both one-sided hypotheses reject at $\alpha = 0.05$.
- **Seeds.** $\{42, 114, 256, 512, 768\}$.

6 HARDWARE AND SOFTWARE ENVIRONMENT

All runs were produced on a single fixed machine; the fingerprint is summarized in Table 1. Wall-clock budgets (60 s / 300 s / 600 s) are interpreted on this machine; the 60 s $\rightarrow$ 300 s plateau (RQ3) is a property of this cost regime.

7 LICENSE AND DATA AVAILABILITY

All datasets used are public DeepLine tabular classification datasets; per-dataset details are listed in Table 2. The file `data/manifest.csv` maps each canonical dataset name to its artifact-relative CSV path, target column, and original DeepLine source URL. If the raw CSV cache is not bundled, users can download the public DeepLine CSVs listed in the manifest and arrange them as `data/deepline_dataset/<dataset>/data.csv` with the corresponding `info.json` target metadata. A local cache can be selected with the `PREPBENCH_DATA_DIR` environment variable.

Table 2: Details of the 56 datasets used for the evaluation: number of instances, number of features, percentage of categorical features, percentage of missing values, and number of labels in the target (number of classes). Dataset names match `data/manifest.csv`, the released result CSVs, and the main-paper figures.

Dataset	# Inst.	# Feat.	% Categ.	% Miss	# Labels
Accident_Casualties	7500	14	0	0	4
adult	7480	15	60	0	2
analcata_data_broadwaymult	285	8	50	1.2	7
analcata_data_germangss	400	6	66.7	0	17
ar4	107	30	0	0	2
bank-full	7368	15	53.3	0	2
baseball	1340	17	5.9	0.1	3
biodeg	1055	42	2.4	0	2
blood-transfusion	748	5	0	0	2
bodyfat	252	15	6.7	0	2
braziltourism	412	9	0	2.6	7
bureau	7125	302	5.3	30.6	2
car	1728	7	100	0	4
chatfield_4	235	13	7.7	0	2
cmc	1473	10	0	0	3
crx	690	16	62.5	0.6	2
data_banknote_authentication	1372	5	0	0	2
dermatology	365	35	0	0.1	6
digggle_table_a2	310	9	11.1	0	2
disclosure_z	662	4	25	0	2
Frogs_MFCCs_family	7195	23	4.3	0	4
Frogs_MFCCs_gen	7195	23	4.3	0	8
Frogs_MFCCs_spec	7195	23	4.3	0	10
glass	214	10	0	0	6
haberman	306	4	0	0	2
home_credit	7250	343	0	19.3	2
HTRU_2	7000	9	0	0	2
imagesegmentation	2310	20	5	0	7
IndianLiverPatientDatasetILPD	583	11	0	0.1	2
Iris	150	5	20	0	3
kc3	458	40	0	0	2
kidney	76	7	42.9	0	2
LendingClubIssuedLoans	7370	53	22.6	3	12
magic04	7450	11	9.1	0	2
mammographic_masses	961	6	0	2.8	2
movement_libras	360	91	0	0	15
no2	500	8	12.5	0	2
plasma_retinol	315	14	28.6	0	2
pm10	500	8	12.5	0	2
schizo	340	14	21.4	17.5	2
Skin_NonSkin	7122	4	0	0	2
socmob	1156	6	83.3	0	2
solar-flare	1066	13	23.1	0	6
test_bng_cmc	55296	10	0	0	3
test_breast	699	10	10.2	0	2
test_credit	1000	21	66.7	0	2
test_eucalyptus	736	20	30	3	5
test_ilpd	583	11	9.1	0	2
test_irish	500	6	66.7	1.1	2
test_phoneme	5404	6	0	0	2
The_broken_machine	7490	59	9.4	2	2
triazines	186	61	1.6	0	2
veteran	137	8	12.5	0	2
weatherAUS	7357	25	24	9.1	2
wilt	4839	6	16.7	0	2
Wine_classification	178	14	0	0	3